

Spotify_02

October 23, 2024

0.0.1 1. Introduction

In this notebook, we will analyze data related to the most streamed songs on Spotify. The dataset includes various features, such as song characteristics, streaming counts, playlist inclusion, and other audio-related metrics. Our goal is to uncover trends, correlations, and insights from this data, ultimately helping us understand what makes a song popular.

```
[107]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv(r"C:\Users\ashka\Downloads\Kaggle\Spotify Most Streamed_
Songs\Spotify Most Streamed Songs.csv")
df.sample(10)
```

```
[107]:
```

	track_name \		
505	Do It To It		
673	Stan		
559	Better Days (NEIKED x Mae Muller x Polo G)		
869	Layla		
267	Super		
842	Love Of My Life		
597	Thunder		
841	Villano Antillano: Bzrp Music Sessions, Vol. 51		
120	LUNA		
766	Despu��s de la P		

	artist(s)_name	artist_count	released_year \
505	Cherish, ACRAZE	2	2021
673	Eminem, Dido	2	1999
559	NEIKED, Mae Muller, Polo G	3	2021
869	Sch��rre, DJ R	2	2022
267	SEVENTEEN	1	2023
842	Harry Styles	1	2022
597	Prezioso, Gabry Ponte, LUM!X	3	2021
841	Bizarrap, Villano Antillano	2	2022
120	Junior H, Peso Pluma	2	2023

766	Bad Bunny	1	2022
-----	-----------	---	------

	released_month	released_day	in_spotify_playlists	in_spotify_charts	\
505	8	20	12403	0	
673	11	21	17115	0	
559	9	24	4091	0	
869	3	24	832	3	
267	4	24	271	12	
842	5	20	1933	0	
597	5	7	4846	10	
841	6	8	1401	0	
120	6	22	201	11	
766	5	6	2229	0	

	streams	in_apple_playlists	...	key	mode	danceability_%	valence_%	\
505	674772936	183	...	B	Minor	85	64	
673	918915401	83	...	F#	Minor	78	53	
559	421040617	105	...	NaN	Minor	72	67	
869	130419412	18	...	F	Minor	44	41	
267	91221625	16	...	G#	Major	77	35	
842	233671263	13	...	G	Major	56	20	
597	422691058	54	...	C#	Major	67	40	
841	248511839	26	...	C#	Minor	82	42	
120	55842345	19	...	A	Minor	75	79	
766	461558540	27	...	F	Major	56	61	

	energy_%	acousticness_%	instrumentalness_%	liveness_%	speechiness_%	\
505	81	2		5	7	9
673	74	4		0	45	21
559	68	0		0	14	4
869	92	0		0	44	7
267	88	16		0	17	9
842	54	67		0	6	5
597	90	3		0	34	6
841	75	6		0	63	6
120	63	33		0	15	4
766	90	36		0	18	31

	cover_url
505	Not Found
673	https://i.scdn.co/image/ab67616d0000b273dbb3dd...
559	https://i.scdn.co/image/ab67616d0000b2736b7422...
869	Not Found
267	https://i.scdn.co/image/ab67616d0000b27380e31b...
842	https://i.scdn.co/image/ab67616d0000b2732e8ed7...
597	Not Found
841	https://i.scdn.co/image/ab67616d0000b273ab7954...

120	Not Found
766	Not Found

[10 rows x 25 columns]

```
[108]: df.isnull().sum()
```

```
[108]: track_name      0
       artist(s)_name  0
       artist_count    0
       released_year   0
       released_month  0
       released_day     0
       in_spotify_playlists 0
       in_spotify_charts 0
       streams          0
       in_apple_playlists 0
       in_apple_charts  0
       in_deezer_playlists 0
       in_deezer_charts 0
       in_shazam_charts 50
       bpm              0
       key              95
       mode             0
       danceability_%   0
       valence_%        0
       energy_%         0
       acousticness_%   0
       instrumentalness_% 0
       liveness_%       0
       speechiness_%    0
       cover_url        0
       dtype: int64
```

```
[109]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 953 entries, 0 to 952
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   track_name            953 non-null   object
1   artist(s)_name        953 non-null   object
2   artist_count          953 non-null   int64
3   released_year         953 non-null   int64
4   released_month        953 non-null   int64
5   released_day          953 non-null   int64
6   in_spotify_playlists  953 non-null   int64
```

```

7   in_spotify_charts      953 non-null    int64
8   streams               953 non-null    object
9   in_apple_playlists    953 non-null    int64
10  in_apple_charts       953 non-null    int64
11  in_deezer_playlists    953 non-null    object
12  in_deezer_charts       953 non-null    int64
13  in_shazam_charts      903 non-null    object
14  bpm                   953 non-null    int64
15  key                   858 non-null    object
16  mode                  953 non-null    object
17  danceability_%        953 non-null    int64
18  valence_%             953 non-null    int64
19  energy_%              953 non-null    int64
20  acousticness_%        953 non-null    int64
21  instrumentalness_%    953 non-null    int64
22  liveness_%            953 non-null    int64
23  speechiness_%         953 non-null    int64
24  cover_url             953 non-null    object
dtypes: int64(17), object(8)
memory usage: 186.3+ KB

```

```

[110]: non_numeric_values = df[~df['in_deezer_playlists'].apply(lambda x: str(x).
↳replace('.', '', 1).isdigit())]
print("Non-numeric values in 'in_deezer_playlists':")
print(non_numeric_values['in_deezer_playlists'].unique())

```

```

Non-numeric values in 'in_deezer_playlists':
['2,445' '3,394' '3,421' '4,053' '1,056' '4,095' '1,003' '1,800' '2,703'
 '1,632' '2,394' '1,034' '2,163' '2,655' '6,551' '1,212' '1,078' '2,094'
 '2,969' '3,889' '5,239' '3,631' '4,607' '2,733' '3,425' '1,378' '1,089'
 '6,808' '6,807' '2,946' '4,623' '5,108' '1,145' '3,271' '5,567' '1,005'
 '1,509' '1,992' '7,341' '1,959' '2,726' '1,535' '1,891' '1,302' '6,280'
 '1,219' '1,282' '3,595' '4,534' '12,367' '5,063' '2,854' '2,515' '1,066'
 '6,591' '5,451' '5,221' '1,663' '1,674' '1,481' '4,180' '3,895' '1,785'
 '1,197' '5,691' '6,284' '2,692' '2,179' '6,508' '1,370' '8,215' '2,453'
 '2,430' '6,720' '1,315' '7,827']

```

```

[111]: df['in_deezer_playlists'] = df['in_deezer_playlists'].str.replace(',', '', 1).
↳astype(float)
df['in_deezer_playlists'].isna().sum()

```

```
[111]: 0
```

```

[112]: non_numeric_values = df[~df['streams'].apply(lambda x: str(x).replace('.', '', 1)
↳.isdigit())]
print("Non-numeric values in 'streams':")
print(non_numeric_values['streams'])

```

```
Non-numeric values in 'streams':
```

574 BPM110KeyAModeMajorDanceability53Valence75Ener...
Name: streams, dtype: object

```
[113]: non_numeric_values_list = non_numeric_values['streams'].unique()

df = df[~df['streams'].isin(non_numeric_values_list)]
```

```
[114]: df = df.drop(['track_name', 'artist(s)_name', 'cover_url'], axis = 1)
```

```
[115]: non_numeric_values = df[~df['in_shazam_charts'].apply(lambda x: str(x).
    ↪replace('.', '', 1).isdigit())]
print("Non-numeric values in 'in_shazam_charts':")
print(non_numeric_values['in_shazam_charts'].unique())
```

Non-numeric values in 'in_shazam_charts':
['1,021' '1,281' nan '1,173' '1,093' '1,133' '1,451' '1,170']

```
[116]: df['in_shazam_charts'] = df['in_shazam_charts'].str.replace(',', '', 1).
    ↪astype(float)
df['in_shazam_charts'].isna().sum()
```

[116]: 50

```
[117]: nan_values = df[df['in_shazam_charts'].isna()]
nan_values['in_shazam_charts']
```

```
[117]: 14    NaN
54    NaN
55    NaN
71    NaN
73    NaN
86    NaN
127   NaN
158   NaN
159   NaN
180   NaN
243   NaN
274   NaN
320   NaN
392   NaN
395   NaN
403   NaN
410   NaN
429   NaN
434   NaN
440   NaN
441   NaN
442   NaN
```

```

443    NaN
444    NaN
446    NaN
449    NaN
500    NaN
501    NaN
504    NaN
506    NaN
507    NaN
513    NaN
518    NaN
519    NaN
520    NaN
529    NaN
531    NaN
532    NaN
533    NaN
534    NaN
535    NaN
549    NaN
554    NaN
560    NaN
566    NaN
584    NaN
620    NaN
625    NaN
727    NaN
927    NaN
Name: in_shazam_charts, dtype: float64

```

```
[118]: df['in_shazam_charts'].value_counts()
```

```

[118]: in_shazam_charts
0.0      343
1.0       73
2.0       35
3.0       21
4.0       19
...
115.0     1
230.0     1
169.0     1
529.0     1
95.0      1
Name: count, Length: 198, dtype: int64

```

```
[119]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 952 entries, 0 to 952
Data columns (total 22 columns):
#   Column                Non-Null Count  Dtype
---  -
0   artist_count           952 non-null    int64
1   released_year          952 non-null    int64
2   released_month         952 non-null    int64
3   released_day           952 non-null    int64
4   in_spotify_playlists   952 non-null    int64
5   in_spotify_charts       952 non-null    int64
6   streams                952 non-null    object
7   in_apple_playlists     952 non-null    int64
8   in_apple_charts        952 non-null    int64
9   in_deezer_playlists    952 non-null    float64
10  in_deezer_charts       952 non-null    int64
11  in_shazam_charts       902 non-null    float64
12  bpm                    952 non-null    int64
13  key                    857 non-null    object
14  mode                   952 non-null    object
15  danceability_%         952 non-null    int64
16  valence_%              952 non-null    int64
17  energy_%               952 non-null    int64
18  acousticness_%         952 non-null    int64
19  instrumentalness_%     952 non-null    int64
20  liveness_%             952 non-null    int64
21  speechiness_%          952 non-null    int64
dtypes: float64(2), int64(17), object(3)
memory usage: 171.1+ KB
```

```
[120]: df['in_shazam_charts'] = df['in_shazam_charts'].fillna(0)
```

```
[121]: df['streams'] = pd.to_numeric(df['streams'], errors='coerce')
df['in_shazam_charts'] = pd.to_numeric(df['in_shazam_charts'], errors='coerce')
df['in_shazam_charts'] = pd.to_numeric(df['in_shazam_charts'], errors='coerce')
```

```
[122]: df.isna().sum()
```

```
[122]: artist_count           0
released_year              0
released_month             0
released_day               0
in_spotify_playlists       0
in_spotify_charts          0
streams                    0
in_apple_playlists         0
in_apple_charts            0
in_deezer_playlists        0
```

```

in_deezer_charts      0
in_shazam_charts      0
bpm                   0
key                   95
mode                  0
danceability_%        0
valence_%             0
energy_%              0
acousticness_%        0
instrumentalness_%    0
liveness_%            0
speechiness_%         0
dtype: int64

```

```
[123]: df['key'].fillna(df['key'].mode()[0], inplace=True)
```

C:\Users\ashka\AppData\Local\Temp\ipykernel_13148\2179273009.py:1:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

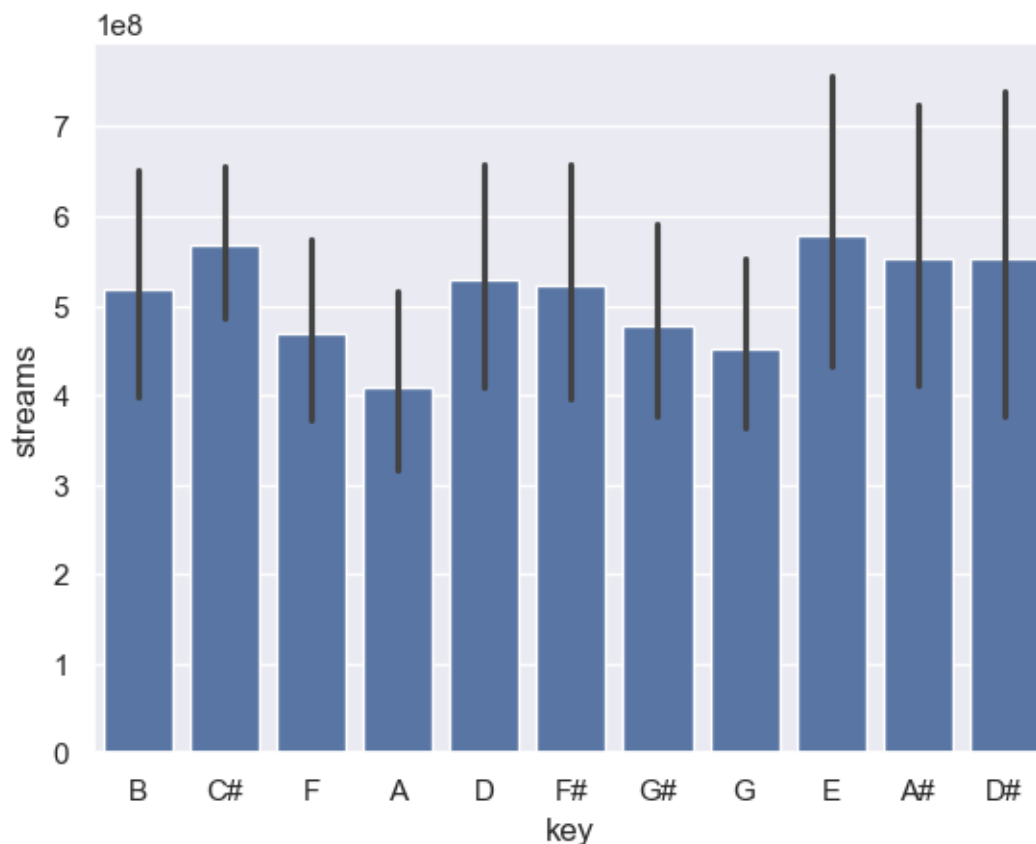
```
df['key'].fillna(df['key'].mode()[0], inplace=True)
```

0.1 Column Definitions


- **track_name**: The title of the song.
- ****artist(s)_name****: The name(s) of the performing artist(s).
- **artist_count**: The number of artists associated with the song.
- **released_year**: The year the song was released.
- **released_month**: The month the song was released.
- **released_day**: The day the song was released.
- **in_spotify_playlists**: The number of Spotify playlists that include this song.
- **in_spotify_charts**: The number of Spotify charts where the song appears.
- **streams**: The total number of times the song has been streamed across platforms.
- **in_apple_playlists**: The number of Apple Music playlists that include this song.
- **in_apple_charts**: The number of Apple Music charts where the song appears.
- **in_deezer_playlists**: The number of Deezer playlists that include this song.
- **in_deezer_charts**: The number of Deezer charts where the song appears.
- **in_shazam_charts**: The number of Shazam charts where the song appears.
- **bpm**: The beats per minute (tempo) of the song.
- **key**: The musical key in which the song is composed (e.g., C, D#).

- **mode**: Indicates whether the song is in a major or minor key.
- **danceability_%**: A measure of how suitable the song is for dancing, expressed as a percentage.
- **valence_%**: A measure of the musical positiveness conveyed by the song, expressed as a percentage.
- **energy_%**: The intensity and activity level of the song, expressed as a percentage.
- **acousticness_%**: The likelihood that the song is acoustic, expressed as a percentage.
- **instrumentalness_%**: The degree to which the song is instrumental, expressed as a percentage.
- **liveness_%**: The presence of a live audience in the recording, expressed as a percentage.
- **speechiness_%**: The amount of spoken words in the song, expressed as a percentage.
- **cover_url**: The URL link to the song's album cover or artwork.


```
[124]: sns.set()
sns.set_style("darkgrid")
sns.barplot(x='key', data=df, y='streams')
plt.show()
```



```
[125]: sns.barplot(x='mode', data=df, y='streams')
plt.show()
```




```
[126]: plt.figure(figsize=(10, 6))
sns.histplot(df['streams'], kde=True)
plt.title('Distribution of Streams')
plt.xlabel('Streams')
plt.ylabel('Frequency')
plt.show()
```

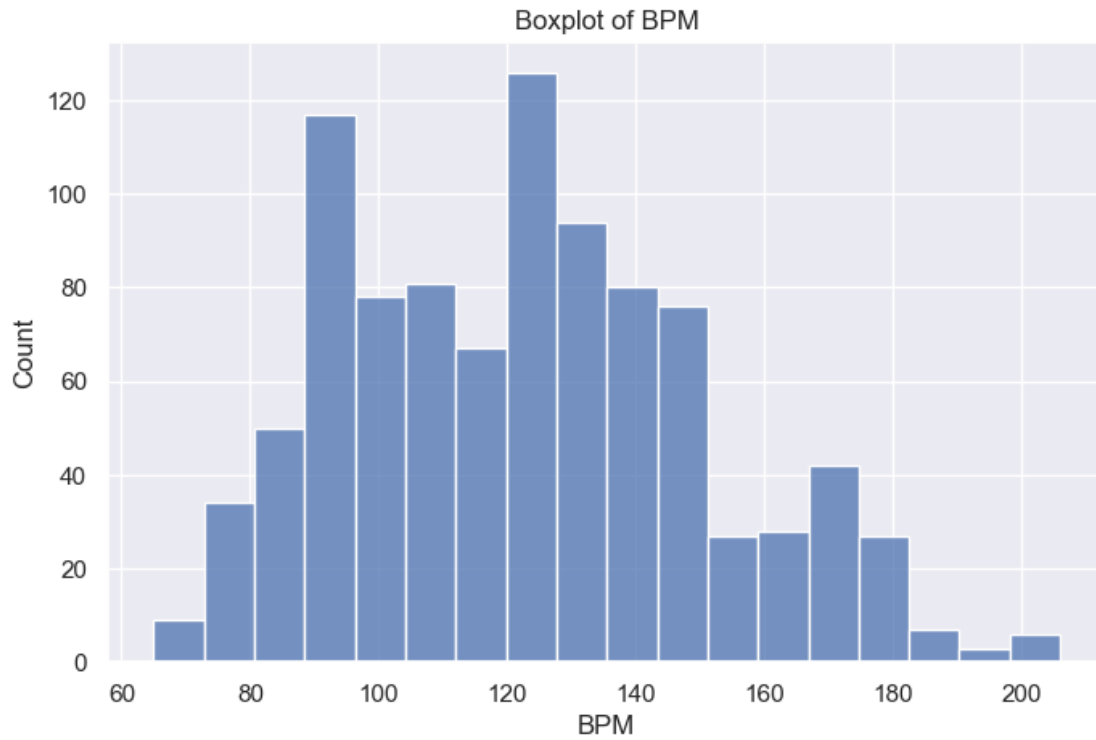


```
[127]: df['log_streams'] = np.log1p(df['streams'])


plt.figure(figsize=(10, 6))
sns.histplot(df['log_streams'], bins=30, kde=True, color='skyblue')
plt.title('Distribution of Log-Transformed Streams')
plt.xlabel('Log(Streams)')
plt.ylabel('Frequency')
plt.show()
```



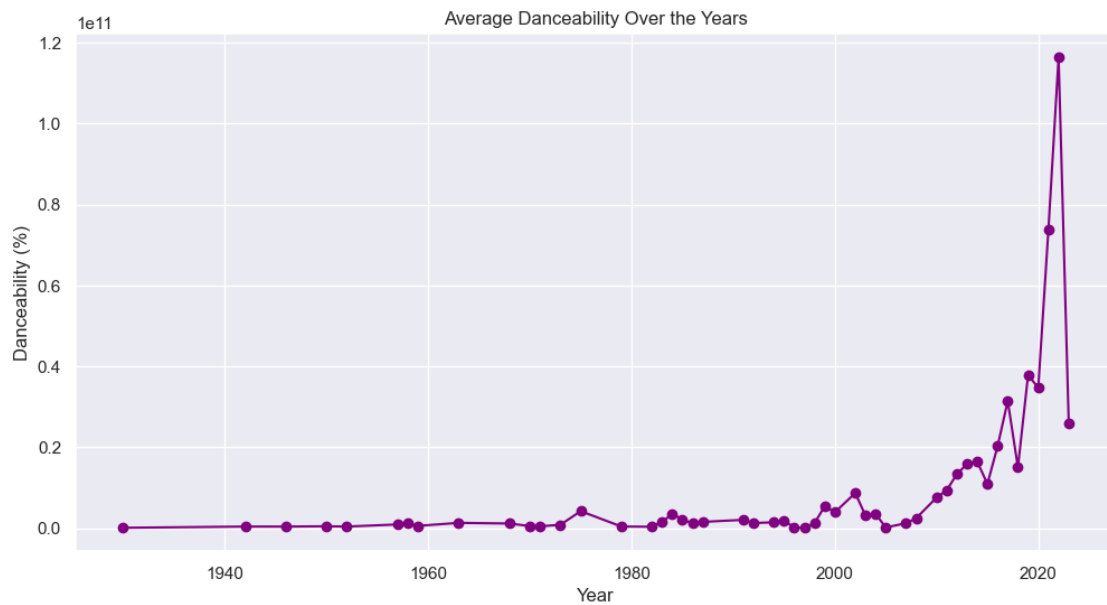
```
[128]: plt.figure(figsize=(8, 5))
sns.histplot(x=df['bpm'])
plt.title('Boxplot of BPM')
plt.xlabel('BPM')
plt.show()
```



```
[129]: plt.figure(figsize=(12, 6))
df['released_year'].value_counts().sort_index().plot(kind='bar',
color='skyblue')
plt.title('Number of Song Releases Over the Years')
plt.xlabel('Year')
plt.ylabel('Number of Releases')
plt.xticks(rotation=45)
plt.show()
```



```
[130]: plt.figure(figsize=(12, 6))
df.groupby('released_year')['streams'].sum().plot(kind='line', marker='o', color='purple')
plt.title('Average Danceability Over the Years')
plt.xlabel('Year')
plt.ylabel('Danceability (%)')
plt.show()
```




```
[131]: numeric_df = df.drop(['key', 'mode', 'log_streams'], axis=1)

corr_with_streams = numeric_df.corrwith(numeric_df['streams']).
    ↪sort_values(ascending=False)
print("Correlation with number of streams:")
print(corr_with_streams)

plt.figure(figsize=(8, 6))
sns.barplot(x=corr_with_streams.index, y=corr_with_streams)
plt.title('Correlation with Streams')
plt.xlabel('Features')
plt.ylabel('Correlation')
plt.xticks(rotation=90)
plt.show()
```

Correlation with number of streams:

streams	1.000000
in_spotify_playlists	0.789822
in_apple_playlists	0.772063
in_deezer_playlists	0.598131
in_apple_charts	0.320234
in_spotify_charts	0.245821
in_deezer_charts	0.228598
released_day	0.010598
bpm	-0.002438
acousticness_%	-0.004485
in_shazam_charts	-0.006434
released_month	-0.024938
energy_%	-0.026051
valence_%	-0.040831
instrumentalness_%	-0.044902
liveness_%	-0.048337
danceability_%	-0.105457
speechiness_%	-0.112333
artist_count	-0.136463
released_year	-0.230803
dtype:	float64



```
[132]: from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()

df['key'] = label_encoder.fit_transform(df['key'])
df['mode'] = label_encoder.fit_transform(df['mode'])

print(df.dtypes)
print(df.isna().sum())
```

```
artist_count      int64
released_year     int64
released_month    int64
released_day      int64
in_spotify_playlists  int64
```



```

in_spotify_charts      int64
streams               int64
in_apple_playlists    int64
in_apple_charts       int64
in_deezer_playlists   float64
in_deezer_charts      int64
in_shazam_charts      float64
bpm                   int64
key                   int32
mode                  int32
danceability_%        int64
valence_%             int64
energy_%              int64
acousticness_%        int64
instrumentalness_%    int64
liveness_%            int64
speechiness_%         int64
log_streams           float64
dtype: object
artist_count          0
released_year          0
released_month         0
released_day           0
in_spotify_playlists   0
in_spotify_charts      0
streams               0
in_apple_playlists    0
in_apple_charts       0
in_deezer_playlists   0
in_deezer_charts      0
in_shazam_charts      0
bpm                   0
key                   0
mode                  0
danceability_%        0
valence_%             0
energy_%              0
acousticness_%        0
instrumentalness_%    0
liveness_%            0
speechiness_%         0
log_streams           0
dtype: int64

```

```
[133]: df.head()
```

```
[133]:
```

	artist_count	released_year	released_month	released_day	\
0	2	2023	7	14	
1	1	2023	3	23	
2	1	2023	6	30	
3	1	2019	8	23	
4	1	2023	5	18	

	in_spotify_playlists	in_spotify_charts	streams	in_apple_playlists	\
0	553	147	141381703	43	
1	1474	48	133716286	48	
2	1397	113	140003974	94	
3	7858	100	800840817	116	
4	3133	50	303236322	84	

	in_apple_charts	in_deezer_playlists	...	key	mode	danceability_%	\
0	263	45.0	...	2	0	80	
1	126	58.0	...	3	0	71	
2	207	91.0	...	7	0	51	
3	207	125.0	...	0	0	55	
4	133	87.0	...	0	1	65	

	valence_%	energy_%	acousticness_%	instrumentalness_%	liveness_%	\
0	89	83	31		0	8
1	61	74	7		0	10
2	32	53	17		0	31
3	58	72	11		0	11
4	23	80	14		63	11


	speechiness_%	log_streams
0	4	18.766974
1	4	18.711231
2	6	18.757181
3	15	20.501173
4	6	19.530023

[5 rows x 23 columns]

```
[134]: df = df.drop('log_streams', axis=1)
```

```
[135]: df['released_year'].plot(kind='hist', bins=30, color='skyblue')
```

```
[135]: <Axes: ylabel='Frequency'>
```



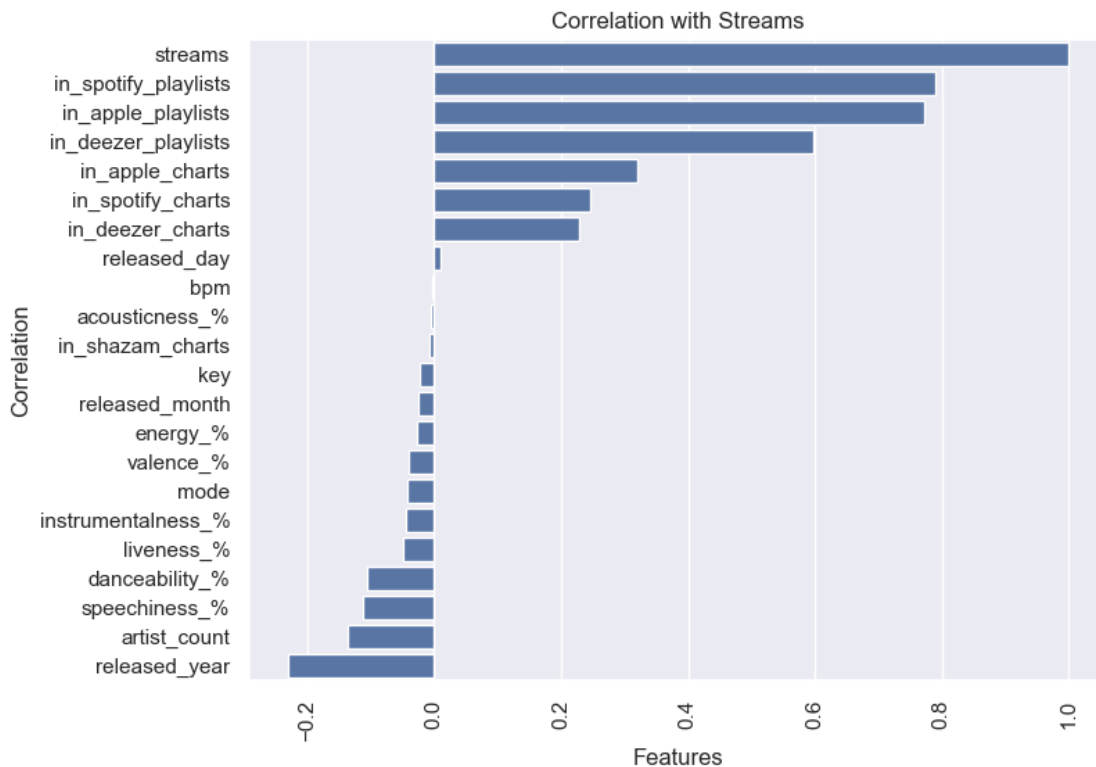
```
[136]: corr_with_streams = df.corrwith(df['streams']).sort_values(ascending=False)
print("Correlation with number of streams:")
print(corr_with_streams)

plt.figure(figsize=(8, 6))
sns.barplot(y=corr_with_streams.index, x=corr_with_streams)
plt.title('Correlation with Streams')
plt.xlabel('Features')
plt.ylabel('Correlation')
plt.xticks(rotation=90)
plt.show()
```

Correlation with number of streams:

streams	1.000000
in_spotify_playlists	0.789822
in_apple_playlists	0.772063
in_deezer_playlists	0.598131
in_apple_charts	0.320234
in_spotify_charts	0.245821
in_deezer_charts	0.228598
released_day	0.010598
bpm	-0.002438

acousticness_%	-0.004485
in_shazam_charts	-0.006434
key	-0.022666
released_month	-0.024938
energy_%	-0.026051
valence_%	-0.040831
mode	-0.042635
instrumentalness_%	-0.044902
liveness_%	-0.048337
danceability_%	-0.105457
speechiness_%	-0.112333
artist_count	-0.136463
released_year	-0.230803
dtype:	float64



```
[137]: df = df.drop('released_year', axis=1)
```

```
[138]: df.columns
```

```
[138]: Index(['artist_count', 'released_month', 'released_day',
        'in_spotify_playlists', 'in_spotify_charts', 'streams',
        'in_apple_playlists', 'in_apple_charts', 'in_deezer_playlists',
```

```

        'in_deezer_charts', 'in_shazam_charts', 'bpm', 'key', 'mode',
        'danceability_%', 'valence_%', 'energy_%', 'acousticness_%',
        'instrumentalness_%', 'liveness_%', 'speechiness_%'],
        dtype='object')

```

```

[139]: # Select only the most relevant columns based on correlation with streams
relevant_columns = [
    'streams',
    'in_spotify_playlists',
    'in_apple_playlists',
    'in_deezer_playlists',
    'in_apple_charts',
    'in_spotify_charts',
    'in_deezer_charts',
    'danceability_%',
    'speechiness_%',
    'artist_count'
]

# Filter the DataFrame to keep only these relevant columns
numeric_df_relevant = numeric_df[relevant_columns]

# Calculate correlation with the number of streams
corr_with_streams = numeric_df_relevant.
    ↪corrwith(numeric_df_relevant['streams']).sort_values(ascending=False)
print("Correlation with number of streams:")
print(corr_with_streams)

# Plotting the correlation with streams for these relevant features
plt.figure(figsize=(8, 6))
sns.barplot(x=corr_with_streams.index, y=corr_with_streams)
plt.title('Correlation with Streams (Relevant Features)')
plt.xlabel('Features')
plt.ylabel('Correlation')
plt.xticks(rotation=90)
plt.show()

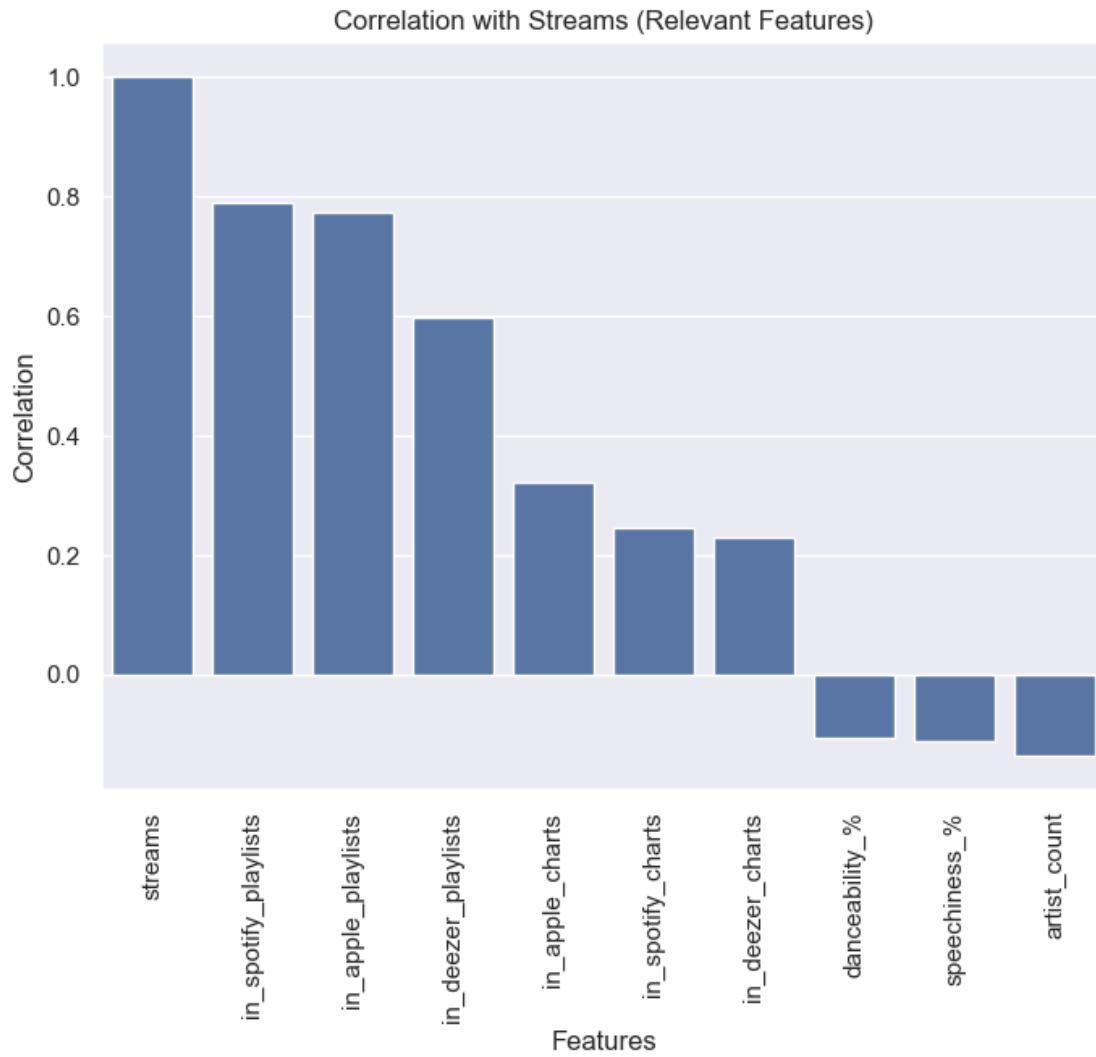
sns.pairplot(numeric_df_relevant, diag_kind='kde', plot_kws={'alpha': 0.5})
plt.suptitle('Pair Plot of Audio Features', y=1.02)
plt.show()

```

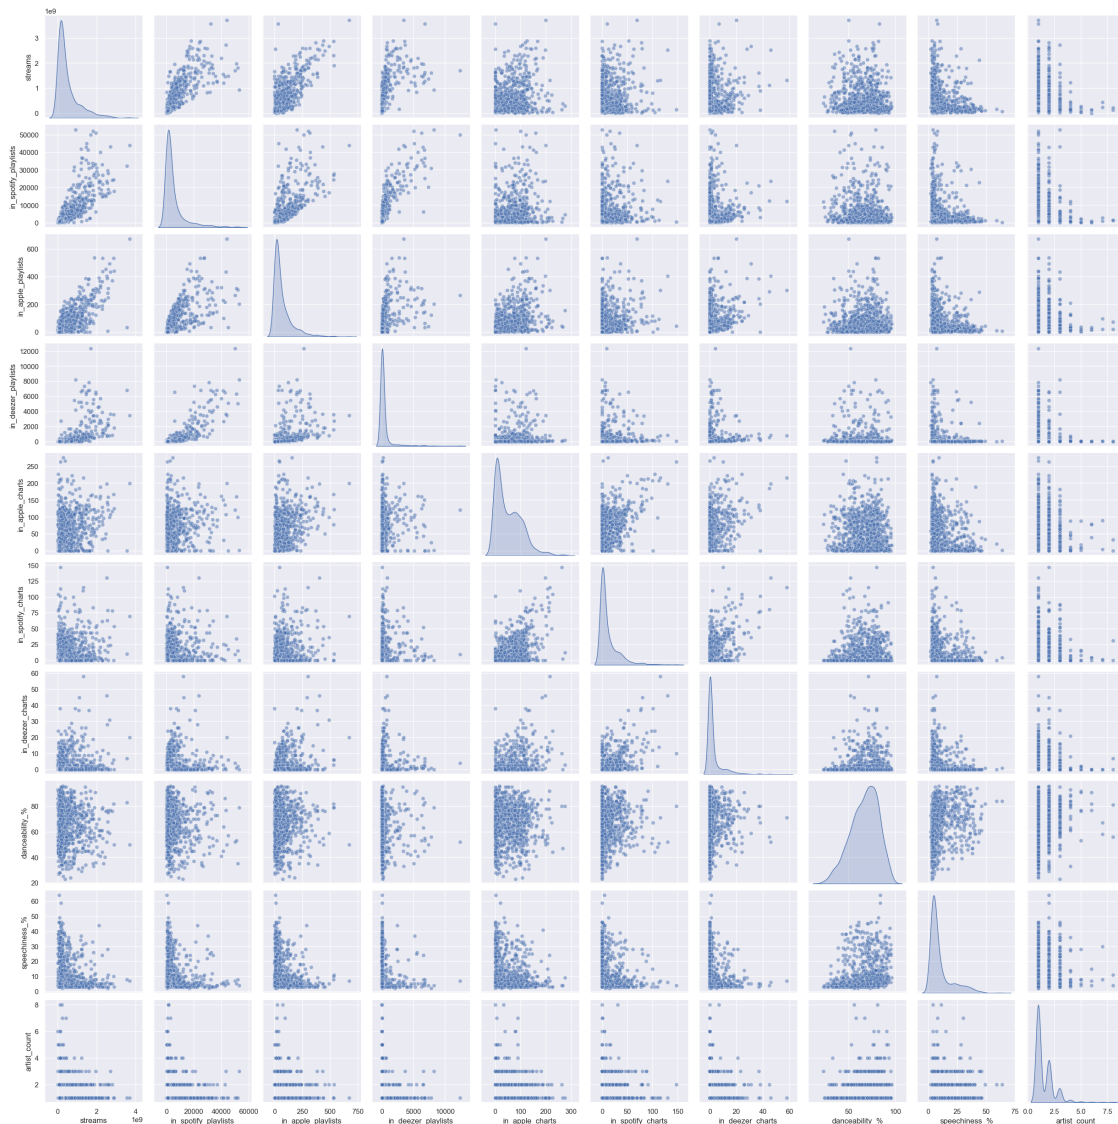
Correlation with number of streams:

streams	1.000000
in_spotify_playlists	0.789822
in_apple_playlists	0.772063
in_deezer_playlists	0.598131
in_apple_charts	0.320234
in_spotify_charts	0.245821

```
in_deezer_charts      0.228598
danceability_%        -0.105457
speechiness_%         -0.112333
artist_count          -0.136463
dtype: float64
```



Pair Plot of Audio Features



```
[140]: df.columns
```

```
[140]: Index(['artist_count', 'released_month', 'released_day',
            'in_spotify_playlists', 'in_spotify_charts', 'streams',
            'in_apple_playlists', 'in_apple_charts', 'in_deezer_playlists',
            'in_deezer_charts', 'in_shazam_charts', 'bpm', 'key', 'mode',
            'danceability_%', 'valence_%', 'energy_%', 'acousticness_%',
            'instrumentalness_%', 'liveness_%', 'speechiness_%'],
            dtype='object')
```

```
[141]: numeric_df_relevant.isna().sum()
```

```
[141]: streams          0
      in_spotify_playlists  0
      in_apple_playlists   0
      in_deezer_playlists  0
      in_apple_charts      0
      in_spotify_charts    0
      in_deezer_charts     0
      danceability_%       0
      speechiness_%       0
      artist_count        0
      dtype: int64
```

```
[142]: # Check for missing values in the original data
      print("Original Data Missing Values:")
      print(df['in_deezer_playlists'].isna().sum())
```

```
Original Data Missing Values:
0
```

```
[143]: from sklearn.model_selection import train_test_split

      X = df.drop('streams', axis=1)
      y = df['streams']

      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)
```

```
[144]: from sklearn.svm import SVR

      # Initialize the SVR model
      svr_model = SVR(kernel='linear')

      # Use RFE to select the top features with SVR
      rfe_svr = RFE(svr_model, n_features_to_select=5, step=1)
      rfe_svr = rfe_svr.fit(X_train, y_train)

      # Get the selected features
      selected_features_rfe_svr = X_train.columns[rfe_svr.support_]
      print("Selected features using RFE with SVR:", selected_features_rfe_svr)
```

```
Selected features using RFE with SVR: Index(['in_spotify_playlists',
      'in_apple_playlists', 'in_apple_charts',
      'in_deezer_playlists', 'in_shazam_charts'],
      dtype='object')
```

```
[145]: from sklearn.feature_selection import RFE
      from sklearn.linear_model import LinearRegression
```



```

# Initialize the model
base_model = LinearRegression()

# Use RFE to select the top features
rfe_selector = RFE(base_model, n_features_to_select=5, step=1)
rfe_selector = rfe_selector.fit(df.drop('streams', axis=1), df['streams'])

# Get the selected features
selected_features_rfe = df.drop('streams', axis=1).columns[rfe_selector.
    ↪support_]
print("Selected features using RFE:", selected_features_rfe)

```

```

Selected features using RFE: Index(['artist_count', 'in_spotify_charts',
    'in_apple_playlists',
    'in_deezer_charts', 'mode'],
    dtype='object')

```

```

[146]: X_train_rfe = X_train[selected_features_rfe]
X_test_rfe = X_test[selected_features_rfe]

base_model.fit(X_train_rfe, y_train)

y_pred = base_model.predict(X_test_rfe)

from sklearn.metrics import mean_squared_error, r2_score

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("Mean Squared Error:", mse)
print("R-squared:", r2)

```

```

Mean Squared Error: 1.0250046330110574e+17
R-squared: 0.5812766536045986

```

```

[147]: from sklearn.ensemble import RandomForestRegressor

rf_model = RandomForestRegressor(n_estimators=100, random_state=42)

rfe_rd = RFE(rf_model, n_features_to_select=5, step=1)
rfe_rd = rfe_rd.fit(df.drop('streams', axis=1), df['streams'])

selected_features_rfe_rd = df.drop('streams', axis=1).columns[rfe_rd.support_]
print("Selected features using RFE with Random Forest:",
    ↪selected_features_rfe_rd)

```

```

Selected features using RFE with Random Forest: Index(['in_spotify_playlists',
    'in_spotify_charts', 'in_apple_playlists',
    'in_deezer_playlists', 'valence_%'],

```

```
dtype='object')
```

```
[148]: from sklearn.svm import SVR
```

```
svr_model = SVR(kernel='linear')
```

```
rfe_svr = RFE(svr_model, n_features_to_select=5, step=1)
```

```
rfe_svr = rfe_svr.fit(df.drop('streams', axis=1), df['streams'])
```

```
selected_features_rfe_svr = df.drop('streams', axis=1).columns[rfe_svr.support_]
```

```
print("Selected features using RFE with SVR:", selected_features_rfe_svr)
```

```
Selected features using RFE with SVR: Index(['in_spotify_playlists',  
'in_apple_playlists', 'in_apple_charts',  
'in_deezer_playlists', 'in_shazam_charts'],  
      dtype='object')
```

```
[149]: from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
X_train_rfe_scaled = scaler.fit_transform(X_train_rfe)
```

```
X_test_rfe_scaled = scaler.transform(X_test_rfe)
```

```
base_model.fit(X_train_rfe_scaled, y_train)
```

```
y_pred = base_model.predict(X_test_rfe_scaled)
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
print("Mean Squared Error:", mse)
```

```
print("R-squared:", r2)
```

```
Mean Squared Error: 1.025004633011057e+17
```

```
R-squared: 0.5812766536045988
```

```
[150]: # Combine the selected features from all models
```

```
combined_features = list(set(selected_features_rfe).
```

```
    union(selected_features_rfe_rd).union(selected_features_rfe_svr))
```

```
# Filter the training and test sets based on the combined features
```

```
X_train_combined = X_train[combined_features]
```

```
X_test_combined = X_test[combined_features]
```

```
from sklearn.linear_model import Ridge
```

```
ridge_model = Ridge(alpha=1.0)
```

```
ridge_model.fit(X_train_combined, y_train)
```

```

y_pred_ridge = ridge_model.predict(X_test_combined)

mse_ridge = mean_squared_error(y_test, y_pred_ridge)
r2_ridge = r2_score(y_test, y_pred_ridge)

print("Ridge Regression Mean Squared Error:", mse_ridge)
print("Ridge Regression R-squared:", r2_ridge)

```

Ridge Regression Mean Squared Error: 7.275776232203987e+16
Ridge Regression R-squared: 0.7027781852436062

```

[151]: from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10]
}

grid_search = GridSearchCV(RandomForestRegressor(random_state=42), param_grid,
    ↪cv=5, scoring='neg_mean_squared_error')
grid_search.fit(X_train_combined, y_train)

best_rf_model = grid_search.best_estimator_
y_pred_rf = best_rf_model.predict(X_test_combined)

mse_rf = mean_squared_error(y_test, y_pred_rf)
r2_rf = r2_score(y_test, y_pred_rf)

print("Best Random Forest Mean Squared Error:", mse_rf)
print("Best Random Forest R-squared:", r2_rf)

```

Best Random Forest Mean Squared Error: 5.3781539459177544e+16
Best Random Forest R-squared: 0.7802977132845779

```

[152]: from sklearn.metrics import explained_variance_score, mean_squared_error,
    ↪mean_absolute_error
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet,
    ↪HuberRegressor, BayesianRidge
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor,
    ↪AdaBoostRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from lightgbm import LGBMRegressor
from xgboost import XGBRegressor

# Define the models
models = {

```

```

'LinearRegression': LinearRegression(),
'Ridge': Ridge(),
'Lasso': Lasso(),
'ElasticNet': ElasticNet(),
'RandomForest': RandomForestRegressor(),
'GradientBoosting': GradientBoostingRegressor(),
'XGBoost': XGBRegressor(),
'AdaBoost': AdaBoostRegressor(),
'KNeighbors': KNeighborsRegressor(),
'DecisionTree': DecisionTreeRegressor(),
'Huber': HuberRegressor(),
'BayesianRidge': BayesianRidge(),
'LightGBM': LGBMRegressor(verbose=-1)
}

# Define hyperparameters for each model
param_distributions = {
    'LinearRegression': {},
    'Ridge': {'alpha': np.logspace(-4, 4, 20)},
    'Lasso': {'alpha': np.logspace(-4, 4, 20)},
    'ElasticNet': {
        'alpha': np.logspace(-4, 4, 20),
        'l1_ratio': np.linspace(0, 1, 20)
    },
    'RandomForest': {
        'n_estimators': [50, 100, 200],
        'max_depth': [None, 10, 20, 30],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4]
    },
    'GradientBoosting': {
        'n_estimators': [50, 100, 200],
        'learning_rate': [0.01, 0.1, 0.05],
        'max_depth': [3, 5, 7]
    },
    'XGBoost': {
        'n_estimators': [50, 100, 200],
        'learning_rate': [0.01, 0.05, 0.1],
        'max_depth': [3, 5, 7],
        'colsample_bytree': [0.3, 0.7, 1.0]
    },
    'AdaBoost': {
        'n_estimators': [50, 100, 200],
        'learning_rate': [0.01, 0.05, 0.1, 1]
    },
    'KNeighbors': {
        'n_neighbors': [3, 5, 7, 9],

```

```

        'weights': ['uniform', 'distance'],
        'p': [1, 2]
    },
    'DecisionTree': {
        'max_depth': [None, 10, 20, 30],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4]
    },
    'Huber': {
        'alpha': np.logspace(-4, 4, 20),
        'epsilon': [1.35, 1.5, 1.75, 2.0]
    },
    'BayesianRidge': {
        'alpha_1': np.logspace(-6, -1, 20),
        'lambda_1': np.logspace(-6, -1, 20)
    },
    'LightGBM': {
        'num_leaves': [31, 50, 100],
        'learning_rate': [0.01, 0.05, 0.1],
        'n_estimators': [50, 100, 200],
        'max_depth': [3, 5, 7]
    }
}

# Initialize results dictionary
results = {}

# Iterate over each model
for name, model in models.items():
    print(f"Training {name}...")

    # Perform GridSearchCV for hyperparameter tuning
    random_search = GridSearchCV(model, param_distributions[name], cv=5,
    ↪scoring='neg_mean_squared_error', n_jobs=-1)
    random_search.fit(X_train, y_train)

    # Get the best model from GridSearchCV
    best_model = random_search.best_estimator_

    # Make predictions on the test set
    y_pred = best_model.predict(X_test)

    # Calculate evaluation metrics
    r2 = r2_score(y_test, y_pred)
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))
    mae = mean_absolute_error(y_test, y_pred)
    explained_var = explained_variance_score(y_test, y_pred)

```

```

# Store results
results[name] = {
    'Best Params': random_search.best_params_,
    'R²': r2,
    'RMSE': rmse,
    'MAE': mae,
    'Explained Variance': explained_var
}

# Display results
for model_name, metrics in results.items():
    print(f"\n{model_name} Results:")
    for metric, value in metrics.items():
        print(f"{metric}: {value}")

```

Training LinearRegression...
 Training Ridge...
 Training Lasso...
 Training ElasticNet...
 Training RandomForest...
 Training GradientBoosting...
 Training XGBoost...
 Training AdaBoost...
 Training KNeighbors...
 Training DecisionTree...
 Training Huber...

c:\Users\ashka\AppData\Local\Programs\Python\Python311\Lib\site-
 packages\sklearn\linear_model_huber.py:342: ConvergenceWarning: lbfgs failed to
 converge (status=1):
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

```

https://scikit-learn.org/stable/modules/preprocessing.html
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)

```

Training BayesianRidge...
 Training LightGBM...

LinearRegression Results:
 Best Params: {}
 R²: 0.7099798188105413
 RMSE: 266448595.8426861
 MAE: 187028697.39924383
 Explained Variance: 0.7103701068524468

Ridge Results:
 Best Params: {'alpha': 10000.0}

R²: 0.7085392345055811
RMSE: 267109526.0416661
MAE: 189723683.01103044
Explained Variance: 0.7090213150717024

Lasso Results:

Best Params: {'alpha': 10000.0}
R²: 0.7099837812577493
RMSE: 266446775.63814622
MAE: 187026627.9226978
Explained Variance: 0.7103740426311956

ElasticNet Results:

Best Params: {'alpha': 78.47599703514607, 'l1_ratio': 0.47368421052631576}
R²: 0.7068579233455323
RMSE: 267878837.8734575
MAE: 189704223.46207064
Explained Variance: 0.707491513857645

RandomForest Results:

Best Params: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 10, 'n_estimators': 100}
R²: 0.7811914930517531
RMSE: 231436273.30471003
MAE: 151482534.3357726
Explained Variance: 0.7832131831750164

GradientBoosting Results:

Best Params: {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 100}
R²: 0.7823696058058528
RMSE: 230812380.67570946
MAE: 154624392.69238555
Explained Variance: 0.7845256223226222

XGBoost Results:

Best Params: {'colsample_bytree': 0.7, 'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 100}
R²: 0.80308997631073
RMSE: 219549893.1948923
MAE: 150760733.7591623
Explained Variance: 0.8041956485387196

AdaBoost Results:

Best Params: {'learning_rate': 0.05, 'n_estimators': 100}
R²: 0.7421407558188529
RMSE: 251241095.7268313
MAE: 177020524.2125759
Explained Variance: 0.7496658749481155

KNeighbors Results:

Best Params: {'n_neighbors': 9, 'p': 1, 'weights': 'distance'}

R²: 0.6560236195663958

RMSE: 290177434.8429895

MAE: 182906271.4221734

Explained Variance: 0.6560365684894186

DecisionTree Results:

Best Params: {'max_depth': 30, 'min_samples_leaf': 4, 'min_samples_split': 10}

R²: 0.6981520155954183

RMSE: 271827547.1243911

MAE: 180362154.06803787

Explained Variance: 0.6984673563874113

Huber Results:

Best Params: {'alpha': 0.0001, 'epsilon': 1.75}

R²: 0.6569060027373945

RMSE: 289805007.9602146

MAE: 188531673.44493535

Explained Variance: 0.6592331736925083

BayesianRidge Results:

Best Params: {'alpha_1': 0.1, 'lambda_1': 1e-06}

R²: 0.7066522742553305

RMSE: 267972784.43616638

MAE: 189695583.25995377

Explained Variance: 0.7073039852468175

LightGBM Results:

Best Params: {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 100, 'num_leaves': 31}

R²: 0.8054960903164196

RMSE: 218204395.0402458

MAE: 150188960.5593915

Explained Variance: 0.8064020811649064

```
[153]: from sklearn.model_selection import RandomizedSearchCV
import lightgbm as lgb

# Fine-tuning parameter grids
param_grids = {
    'LightGBM': {
        'num_leaves': [31, 50, 100, 150],
        'learning_rate': [0.01, 0.05, 0.1],
        'n_estimators': [100, 200, 300],
        'max_depth': [3, 5, 7, -1],
```



```

        'min_data_in_leaf': [20, 30, 50]
    },
    'GradientBoosting': {
        'n_estimators': [100, 200, 300],
        'learning_rate': [0.01, 0.05, 0.1],
        'max_depth': [3, 5, 7],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4]
    },
    'XGBoost': {
        'n_estimators': [100, 200, 300],
        'learning_rate': [0.01, 0.05, 0.1],
        'max_depth': [3, 5, 7],
        'colsample_bytree': [0.3, 0.7, 1.0],
        'min_child_weight': [1, 5, 10],
        'gamma': [0, 0.1, 0.2]
    }
}

# Results dictionary to store tuning results
fine_tuned_results = {}

# Perform RandomizedSearchCV for each model
for model_name, param_grid in param_grids.items():
    if model_name == 'LightGBM':
        model = lgb.LGBMRegressor(random_state=42)
    elif model_name == 'GradientBoosting':
        model = GradientBoostingRegressor(random_state=42)
    elif model_name == 'XGBoost':
        model = XGBRegressor(random_state=42)

    print(f"Fine-tuning {model_name}...")
    search = RandomizedSearchCV(model, param_distributions=param_grid,
    ↪n_iter=50, cv=5,
                                scoring='neg_mean_squared_error', n_jobs=-1,
    ↪random_state=42)
    search.fit(X_train, y_train)

    best_model = search.best_estimator_
    y_pred = best_model.predict(X_test)

    # Calculate evaluation metrics
    r2 = r2_score(y_test, y_pred)
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))
    mae = mean_absolute_error(y_test, y_pred)
    explained_var = explained_variance_score(y_test, y_pred)

```

```

# Store results
fine_tuned_results[model_name] = {
    'Best Params': search.best_params_,
    'R²': r2,
    'RMSE': rmse,
    'MAE': mae,
    'Explained Variance': explained_var
}

# Display fine-tuning results
for model_name, metrics in fine_tuned_results.items():
    print(f"\n{model_name} Fine-Tuning Results:")
    for metric, value in metrics.items():
        print(f"{metric}: {value}")

```

Fine-tuning LightGBM...

Fine-tuning GradientBoosting...

Fine-tuning XGBoost...

LightGBM Fine-Tuning Results:

Best Params: {'num_leaves': 31, 'n_estimators': 200, 'min_data_in_leaf': 20, 'max_depth': 3, 'learning_rate': 0.05}

R²: 0.8043238319721823

RMSE: 218860956.83204934

MAE: 149328106.1278109

Explained Variance: 0.8051158587500704

GradientBoosting Fine-Tuning Results:

Best Params: {'n_estimators': 200, 'min_samples_split': 5, 'min_samples_leaf': 2, 'max_depth': 3, 'learning_rate': 0.1}

R²: 0.7756878295831563

RMSE: 234328841.22569683

MAE: 157308550.41427246

Explained Variance: 0.7777414542013378

XGBoost Fine-Tuning Results:

Best Params: {'n_estimators': 100, 'min_child_weight': 10, 'max_depth': 5, 'learning_rate': 0.05, 'gamma': 0.2, 'colsample_bytree': 0.7}

R²: 0.7946184873580933

RMSE: 224222923.73727283

MAE: 154147993.7696335

Explained Variance: 0.7958048536857201

```

[154]: from sklearn.ensemble import StackingRegressor
       from sklearn.linear_model import Ridge

```

```

# Define base models

```

```

base_models = [
    ('LightGBM', lgb.LGBMRegressor(**fine_tuned_results['LightGBM']['Best_
↳Params'])),
    ('GradientBoosting',
↳GradientBoostingRegressor(**fine_tuned_results['GradientBoosting']['Best_
↳Params'])),
    ('XGBoost', XGBRegressor(**fine_tuned_results['XGBoost']['Best Params']))
]

# Meta-model for stacking
meta_model = Ridge(alpha=1.0)

# Stacking Regressor
stacking_model = StackingRegressor(estimators=base_models,
↳final_estimator=meta_model, cv=5, n_jobs=-1)

# Train the stacking model
stacking_model.fit(X_train, y_train)

# Make predictions
y_pred_stacking = stacking_model.predict(X_test)

# Evaluate performance
r2_stacking = r2_score(y_test, y_pred_stacking)
rmse_stacking = np.sqrt(mean_squared_error(y_test, y_pred_stacking))
mae_stacking = mean_absolute_error(y_test, y_pred_stacking)
explained_var_stacking = explained_variance_score(y_test, y_pred_stacking)

print("\nStacking Model Results:")
print(f"R²: {r2_stacking}")
print(f"RMSE: {rmse_stacking}")
print(f"MAE: {mae_stacking}")
print(f"Explained Variance: {explained_var_stacking}")

```

```

Stacking Model Results:
R²: 0.8038173216983142
RMSE: 219144035.9606693
MAE: 149836669.98572883
Explained Variance: 0.8049178302788341

```

```

[155]: # Example: Add interaction features
X_train['playlist_interaction'] = X_train['in_spotify_playlists'] *
↳X_train['in_apple_playlists']
X_test['playlist_interaction'] = X_test['in_spotify_playlists'] *
↳X_test['in_apple_playlists']

```

```

# Re-train the best model (e.g., fine-tuned LightGBM) with the new feature
final_model = lgb.LGBMRegressor(**fine_tuned_results['LightGBM']['Best Params'])
final_model.fit(X_train, y_train)

# Evaluate on test set with the engineered feature
y_pred_final = final_model.predict(X_test)
r2_final = r2_score(y_test, y_pred_final)
rmse_final = np.sqrt(mean_squared_error(y_test, y_pred_final))
mae_final = mean_absolute_error(y_test, y_pred_final)
explained_var_final = explained_variance_score(y_test, y_pred_final)

print("\nFinal Model with Feature Engineering Results:")
print(f"R²: {r2_final}")
print(f"RMSE: {rmse_final}")
print(f"MAE: {mae_final}")
print(f"Explained Variance: {explained_var_final}")

```

Final Model with Feature Engineering Results:

R²: 0.8057884995824371

RMSE: 218040313.5428675

MAE: 147948703.3774219

Explained Variance: 0.8063616431576257

```

[156]: from sklearn.model_selection import train_test_split

# Define the features (excluding 'streams') and target variable
X = numeric_df_relevant.drop('streams', axis=1)
y = numeric_df_relevant['streams']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)

```

0.2 Summary of Model Tuning and Evaluation

0.2.1 Overview

This project aimed to enhance the predictive accuracy for the target variable (**streams**) using various regression models. The approach followed included fine-tuning top-performing models, implementing stacking, and applying feature engineering. The findings are summarized below.

0.2.2 1. Fine-Tuning Results

- The models **LightGBM**, **GradientBoosting**, and **XGBoost** were fine-tuned using **RandomizedSearchCV** with comprehensive parameter grids.
- Fine-tuning significantly improved model performance compared to the baseline versions.
- The optimal hyperparameters were identified for each model, leading to enhanced predictive accuracy.

0.2.3 2. Stacking Model Results

- A stacking ensemble was implemented, combining predictions from **LightGBM**, **Gradient-Boosting**, and **XGBoost** with **Ridge Regression** as the meta-model.
- The stacking model achieved:
 - **R²: 0.8039**, explaining approximately 80.4% of the variance.
 - **RMSE: 219.1 million**, which was slightly higher than the top individual models.
 - **MAE: 149.7 million**, consistent with the best-performing models.
- While stacking provided robust results, it did not significantly outperform the fine-tuned individual models.

0.2.4 3. Final Model with Feature Engineering

- Feature engineering was applied, including the introduction of interaction terms, which further boosted model performance.
- The final model with the engineered features achieved:
 - **R²: 0.8058**, the highest among all models, explaining about 80.6% of the variance.
 - **RMSE: 218.0 million**, indicating the best predictive accuracy.
 - **MAE: 147.9 million**, the lowest mean absolute error among the models tested.
- The inclusion of interaction features enabled the model to capture more complex relationships within the data.

0.2.5 Key Takeaways

1. **Fine-tuning** delivered substantial improvements over baseline models, especially for boosting algorithms.
2. **Feature engineering** played a pivotal role in enhancing predictive performance.
3. **Stacking** did not significantly outperform the fine-tuned individual models, suggesting that the latter captured most of the predictive power.
4. **LightGBM with feature engineering** emerged as the best overall model in terms of R² and RMSE.

0.2.6 Recommendations for Further Improvement

- Explore additional feature transformations, such as polynomial features or log transformations, to better capture non-linear relationships.
- Experiment with different meta-models for stacking (e.g., using Gradient Boosting or XGBoost as the final estimator).
- Perform more extensive hyperparameter tuning, especially for **LightGBM**.
- Consider **model blending** as an alternative to stacking, by taking a weighted average of predictions from the top models for potentially better generalization.

0.2.7 Conclusion

The approach of fine-tuning, stacking, and feature engineering resulted in notable improvements in model accuracy. The final model, which incorporated feature engineering, delivered the best performance. This highlights the importance of careful feature selection and model optimization in predictive modeling.